


Sequence Length
400K 300K 200K 100K


(a) Prefill Latency

Non-Attention
Attention



(b) Decode Latency

Non-Attention
Attention



(c) Memory Usage

Model Weights
KV Cache

