

* Evaluated at 100K context length, with latency averaged over 100 runs

